


OpenCL exercise 5: Volume rendering

Kaicong Sun

Volume rendering



- ▶ Ray goes from origin to the output pixels
- ▶ Values of object (= input data) along the ray are summed up
- ▶ If value is not taken in the middle of a pixel, trilinear interpolation is used (bilinear in 2D-case)
- ▶ Sum of the values is value for output pixel
- ▶ Values outside the input object = 0

Task

- ▶ GPU implementation of 3D volume rendering
 - ▶ Use 3D image object for input data.
 - ▶ Each pixel on the display is a work-item.
- ▶ Profiling code which prints the CPU time / GPU time / memory transfer and speedups.
 - ▶ For memory transfer: Only time for transferring output data

Hints on Host

```
//The inverse view matrix (put camera in world space)  
d_invViewMatrix is a floating points matrix with  
16 elements (defined in render() function).
```

```
//For writing 3D image:  
queue.enqueueWriteImage(d_input, true, origin, region,  
countX * sizeof (float), countX * countY * sizeof (float),  
(void*) h_input, NULL, &copyToDev);
```

```
//For launching the kernel:  
work-group size can also be cl::NullRange in which case  
the OpenCL implementation will determine how to break the  
global work-items into appropriate work-group instances.
```

Hints on Device

// Data type and operations:

OpenCL library has defined float3, min(), max(),
normalize(), dot()

//For getting image size:

float3 boxMax = (float3)get_image_width(d_input),
get_image_height(d_input), get_image_depth(d_input)

// using CLK_FILTER_LINEAR in sampler to have
interpolated value (note for 3D image, coordinates are
float4/int4):

float sample = read_imagef(d_input, sampler,
(float4)(pos, 0)).x