

OpenCL exercise 6: Prefix sum

Kaicong Sun

Prefix sum

- ▶ Prefix sum = all prefix sums for an input vector
- ▶ For input values x_0, x_1, x_2, \dots compute:

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

- ▶ Also can be some other associative binary operation instead of $+$, e.g. min, max, ...


Prefix sum

Host code:

```
1 cl_int sum = h_input[0];
2 h_output[0] = sum;
3 for (std::size_t i = 1; i < h_input.size (); i++) {
4     sum += h_input[i];
5     h_output[i] = sum;
6 }
```


Parallel prefix sum

- Parallel prefix sum:



GPU

- ▶ Problem: Can use only one work group
- ▶ Solution: Work with blocks



- ▶ Launch Kernel1: Do prefix sum per block and write the sum of each block to a temp array until there is only one work group. (For details see another .PDF file)
- ▶ Launch Kernel2: For all blocks except the first: Add `temp2[blockIndex-1]` to all values in the current `temp1` block and then recursively add `temp1[blockIndex-1]` to all values in the current `d_output`.

GPU

- ▶ Task: Implement prefix sum on GPU
 - ▶ Plus usual code for performance measurements
- ▶ Kernel1 should calculate prefix sum blockwise:
 - ▶ Load input data to local memory
 - ▶ Loop over offsets
 - ▶ Write results to global memory
 - ▶ Use one work item per value
 - ▶ Do not forget to add `barrier` calls for synchronization
- ▶ Kernel2 should add the sum of previous blocks to each element of the current block.

Hints on Device

```
//To locate your work group:  
uint grpid = get_group_id(0);
```

```
//Can not use:  
if (li >= offset)  
ldata[li] += ldata[li - offset];  
barrier(CLK_LOCAL_MEM_FENCE);
```

Why?

Solution? \Rightarrow Instead of updating `ldata[li]` directly, using temporary variable and `barrier(CLK_LOCAL_MEM_FENCE)`.